

Package ‘is2’

April 8, 2015

Type Package

Title Iterated smoothing for partially observed Markov processes

Version 0.10-1

Date 2015-04-03

URL <http://is2.r-forge.r-project.org>

Description Inference methods using iterated smoothing for partially-observed Markov processes

Depends R(>= 3.0.0), pomp

Imports stats, graphics, mvtnorm, deSolve, coda

License GPL(>= 2)

LazyData true

Collate generics.R pfilter2.R pfilter2-methods.R is2.R is2-methods.R

Author Dao Nguyen [aut, cre], Edward L. Ionides [aut]

Maintainer Dao Nguyen <nguyenxd@umich.edu>

Repository R-Forge

Repository/R-Forge/Project is2

Repository/R-Forge/Revision 1.0

Repository/R-Forge/DateTimeStamp 2015-04-04 19:23:13

Date/Publication 2015-04-04 19:23:13

R topics documented:

is2	2
is2-methods	5
pfilter2	6
pfilter2-methods	8
Index	10

Description

Iterated smoothing algorithms for estimating the parameters of a partially-observed Markov process.

Usage

```
## S4 method for signature pomp
is2(object, Nis = 1, start, pars, ivps = character(0),
    particles, rw.sd, Np, ic.lag, var.factor, lag,
    cooling.type = c("geometric", "hyperbolic"),
    cooling.fraction, cooling.factor,
    method = c("is2", "ris1", "is1"),
    tol = 1e-17, max.fail = Inf,
    verbose = getOption("verbose"), transform = FALSE, ...)
## S4 method for signature pfiltered2.pomp
is2(object, Nis = 1, Np, tol, ...)
## S4 method for signature is2
is2(object, Nis, start, pars, ivps,
    particles, rw.sd, Np, ic.lag, lag, var.factor,
    cooling.type, cooling.fraction,
    method, tol, transform, ...)
## S4 method for signature is2
continue(object, Nis = 1, ...)
```

Arguments

<code>object</code>	An object of class <code>pomp</code> .
<code>Nis</code>	The number of filtering iterations to perform.
<code>start</code>	named numerical vector; the starting guess of the parameters.
<code>pars</code>	optional character vector naming the ordinary parameters to be estimated. Every parameter named in <code>pars</code> must have a positive random-walk standard deviation specified in <code>rw.sd</code> . Leaving <code>pars</code> unspecified is equivalent to setting it equal to the names of all parameters with a positive value of <code>rw.sd</code> that are not <code>ivps</code> .
<code>ivps</code>	optional character vector naming the initial-value parameters (IVPs) to be estimated. Every parameter named in <code>ivps</code> must have a positive random-walk standard deviation specified in <code>rw.sd</code> . If <code>pars</code> is empty, i.e., only IVPs are to be estimated, see below “Using <code>is2</code> to estimate initial-value parameters only”.
<code>particles</code>	Function of prototype <code>particles(Np, center, sd, ...)</code> which sets up the starting particle matrix by drawing a sample of size <code>Np</code> from the starting particle distribution centered at <code>center</code> and of width <code>sd</code> . If <code>particles</code> is not supplied by the user, the default behavior is to draw the particles from a multivariate normal distribution with mean <code>center</code> and standard deviation <code>sd</code> .

<code>rw.sd</code>	numeric vector with names; the intensity of the random walk to be applied to parameters. The random walk is only applied to parameters named in <code>pars</code> (i.e., not to those named in <code>ivps</code>). The algorithm requires that the random walk be nontrivial, so each element in <code>rw.sd[pars]</code> must be positive. <code>rw.sd</code> is also used to scale the initial-value parameters (via the <code>particles</code> function). Therefore, each element of <code>rw.sd[ivps]</code> must be positive. The following must be satisfied: <code>names(rw.sd)</code> must be a subset of <code>names(start)</code> , <code>rw.sd</code> must be non-negative (zeros are simply ignored), the name of every positive element of <code>rw.sd</code> must be in either <code>pars</code> or <code>ivps</code> .
<code>Np</code>	the number of particles to use in filtering. This may be specified as a single positive integer, in which case the same number of particles will be used at each timestep. Alternatively, if one wishes the number of particles to vary across timestep, one may specify <code>Np</code> either as a vector of positive integers (of length <code>length(time(object, t0=TRUE))</code>) or as a function taking a positive integer argument. In the latter case, <code>Np(k)</code> must be a single positive integer, representing the number of particles to be used at the k-th timestep: <code>Np(0)</code> is the number of particles to use going from <code>timezero(object)</code> to <code>time(object)[1]</code> , <code>Np(1)</code> , from <code>timezero(object)</code> to <code>time(object)[1]</code> , and so on, while when <code>T=length(time(object, t0=TRUE))</code> , <code>Np(T)</code> is the number of particles to sample at the end of the time-series.
<code>ic.lag</code>	a positive integer; the timepoint for fixed-lag smoothing of initial-value parameters. The <code>is2</code> update for initial-value parameters consists of replacing them by their filtering mean at time <code>times[ic.lag]</code> , where <code>times=time(object)</code> . It makes no sense to set <code>ic.lag>length(times)</code> ; if it is so set, <code>ic.lag</code> is set to <code>length(times)</code> with a warning.
<code>var.factor</code>	a positive number; the scaling coefficient relating the width of the starting particle distribution to <code>rw.sd</code> . In particular, the width of the distribution of particles at the start of the first <code>is2</code> iteration will be <code>random.walk.sd*var.factor</code> .
<code>lag</code>	a positive integer; the timepoint for fixed-lag smoothing parameters estimation.
<code>cooling.type</code> , <code>cooling.fraction</code> , <code>cooling.factor</code>	specifications for the cooling schedule, i.e., the manner in which the intensity of the parameter perturbations is reduced with successive filtering iterations. <code>cooling.type</code> specifies the nature of the cooling schedule. When <code>cooling.type="geometric"</code> , on the n-th <code>is2</code> iteration, the relative perturbation intensity is <code>cooling.fraction^(n/50)</code> . When <code>cooling.type="hyperbolic"</code> , on the n-th <code>is2</code> iteration, the relative perturbation intensity is <code>(s+1)/(s+n)</code> , where <code>(s+1)/(s+50)=cooling.fraction</code> . <code>cooling.fraction</code> is the relative magnitude of the parameter perturbations after 50 <code>is2</code> iterations. <code>cooling.factor</code> is now deprecated: to achieve the old behavior, use <code>cooling.type="geometric"</code> and <code>cooling.fraction=(cooling.factor)^50</code> .
<code>method</code>	<code>method</code> sets the update rule used in the algorithm. <code>method="is2"</code> uses the iterated smoothing update rule (Ionides 2015 submitted); <code>method="ris1"</code> uses the reduced iterated smoothing update rule (Doucet 2013, Ionides 2015 submitted); <code>method="is1"</code> uses the reduced iterated smoothing update rule (Doucet 2013);
<code>tol</code>	See the description under pfilter2 .
<code>max.fail</code>	See the description under pfilter2 .

verbose	logical; if TRUE, print progress reports.
transform	logical; if TRUE, optimization is performed on the transformed scale.
...	additional arguments that override the defaults.

Re-running is2 Iterations

To re-run a sequence of is2 iterations, one can use the `is2` method on a `is2` object. By default, the same parameters used for the original is2 run are re-used (except for `weighted`, `tol`, `max.fail`, and `verbose`, the defaults of which are shown above). If one does specify additional arguments, these will override the defaults.

Continuing is2 Iterations

One can resume a series of is2 iterations from where one left off using the `continue` method. A call to `is2` to perform $N_{is}=m$ iterations followed by a call to `continue` to perform $N_{is}=n$ iterations will produce precisely the same effect as a single call to `is2` to perform $N_{is}=m+n$ iterations. By default, all the algorithmic parameters are the same as used in the original call to `is2`. Additional arguments will override the defaults.

Using is2 to estimate initial-value parameters only

One can use is2 fixed-lag smoothing to estimate only initial value parameters (IVPs). In this case, `pars` is left empty and the IVPs to be estimated are named in `ivps`. If `theta` is the current parameter vector, then at each is2 iteration, N_p particles are drawn from a distribution centered at `theta` and with width proportional to `var.factor*rw.sd`, a particle filtering operation is performed, and `theta` is replaced by the filtering mean at `time(object)[ic.lag]`. Note the implication that, when is2 is used in this way on a time series any longer than `ic.lag`, unnecessary work is done. If the time series in `object` is longer than `ic.lag`, consider replacing `object` with `window(object, end=ic.lag)`.

Details

If `particles` is not specified, the default behavior is to draw the particles from a multivariate normal distribution. It is the user's responsibility to ensure that, if the optional `particles` argument is given, that the `particles` function satisfies the following conditions:

`particles` has at least the following arguments: `Np`, `center`, `sd`, and ... `Np` may be assumed to be a positive integer; `center` and `sd` will be named vectors of the same length. Additional arguments may be specified; these will be filled with the elements of the `userdata` slot of the underlying `pomp` object (see [pomp](#)).

`particles` returns a `length(center) x Np` matrix with `rownames` matching the names of `center` and `sd`. Each column represents a distinct particle.

The center of the particle distribution returned by `particles` should be `center`. The width of the particle distribution should vary monotonically with `sd`. In particular, when `sd=0`, the `particles` should return matrices with `Np` identical columns, each given by the parameters specified in `center`.

Author(s)

Dao Nguyen <nguyend at umich dot edu>, Edward L. Ionides <ionides at umich dot edu>

References

- D. Nguyen, E. L. Ionides, A second-order iterated smoothing, Journal of Statistics Society Series B, 2015 (submitted).
- A., Doucet, P. E. Jacob, and S. Rubenthaler, Derivative-free estimation of the score vector and observed information matrix with application to state-space models, Preprint arXiv:1304.5768.
- A. A. King, D. Nguyen, and E. L. Ionides, Statistical inference for partially observed Markov processes via the R package pomp, Journal of Statistical Software, 2015 (just-accepted).

See Also

[is2-methods](#), [pfilter2](#). See the “intro_to_is2” vignette for examples.

is2-methods

Methods of the "is2" class

Description

Methods of the is2 class.

Usage

```
## S4 method for signature is2
logLik(object, ...)
## S4 method for signature is2
conv.rec(object, pars, transform = FALSE, ...)
## S4 method for signature is2List
conv.rec(object, ...)
## S4 method for signature is2
plot(x, y, ...)
## S4 method for signature is2List
plot(x, y, ...)
## S4 method for signature is2
c(x, ..., recursive = FALSE)
## S4 method for signature is2List
c(x, ..., recursive = FALSE)
compare.is2(z)
```

Arguments

object	The is2 object.
pars	Names of parameters.
x	The is2 object.
y, recursive	Ignored.
z	A is2 object or list of is2 objects.
transform	optional logical; should the parameter transformations be applied? See coef for details.
...	Further arguments (either ignored or passed to underlying functions).

Methods

conv.rec `conv.rec(object, pars = NULL)` returns the columns of the convergence-record matrix corresponding to the names in `pars`. By default, all rows are returned.

logLik Returns the value in the `loglik` slot.

c Concatenates `is2` objects into an `is2List`.

plot Plots a series of diagnostic plots.

compare.is2 Deprecated: use `plot` instead.

Author(s)

Dao Nguyen <nguyenxd at umich dot edu>, Edward L. Ionides <ionides at umich dot edu>

See Also

[is2](#), [pfilter2](#)

pfilter2

Particle filter

Description

Run a plain vanilla particle filter. Resampling is performed at each observation.

Usage

```
## S4 method for signature pomp
pfilter2(object, params, Np, tol = 1e-17,
         max.fail = Inf, pred.mean = FALSE, pred.var = FALSE,
         filter.mean = FALSE,
         save.states = FALSE,
         save.params = FALSE, lag=0, seed = NULL,
         verbose = getOption("verbose"), ...)
## S4 method for signature pfilterd2.pomp
pfilter2(object, params, Np, tol, ...)
```

Arguments

<code>object</code>	An object of class <code>pomp</code> or inheriting class <code>pomp</code> .
<code>params</code>	A <code>npars</code> x <code>Np</code> numeric matrix containing the parameters corresponding to the initial state values in <code>xstart</code> . This must have a ‘ <code>rownames</code> ’ attribute. If it desired that all particles should share the same parameter values, one one may supply <code>params</code> as a named numeric vector.

Np	the number of particles to use. This may be specified as a single positive integer, in which case the same number of particles will be used at each timestep. Alternatively, if one wishes the number of particles to vary across timesteps, one may specify Np either as a vector of positive integers (<code>length(time(object,t0=TRUE))</code>) or as a function taking a positive integer argument. In the latter case, Np(k) must be a single positive integer, representing the number of particles to be used at the k-th timestep: Np(0) is the number of particles to use going from <code>timezero(object)</code> to <code>time(object)[1]</code> , Np(1), from <code>timezero(object)</code> to <code>time(object)[1]</code> , and so on, while when <code>T=length(time(object,t0=TRUE))</code> , Np(T) is the number of particles to sample at the end of the time-series. When object is of class <code>is2</code> , this is by default the same number of particles used in the <code>is2</code> iterations.
tol	positive numeric scalar; particles with likelihood less than tol are considered to be “lost”. A filtering failure occurs when, at some time point, all particles are lost. When all particles are lost, the conditional likelihood at that time point is set to tol.
max.fail	integer; the maximum number of filtering failures allowed. If the number of filtering failures exceeds this number, execution will terminate with an error. By default, max.fail is set to infinity, so no error can be triggered.
pred.mean	logical; if TRUE, the prediction means are calculated for the state variables and parameters.
pred.var	logical; if TRUE, the prediction variances are calculated for the state variables and parameters.
filter.mean	logical; if TRUE, the filtering means are calculated for the state variables and parameters.
save.states, save.params	logical. If <code>save.states=TRUE</code> , the state-vector for each particle at each time is saved in the <code>saved.states</code> slot of the returned <code>pfilterd2.pomp</code> object. If <code>save.params=TRUE</code> , the parameter-vector for each particle at each time is saved in the <code>saved.params</code> slot of the returned <code>pfilterd2.pomp</code> object.
lag	positive numeric scalar; use for fixed lag smoothing.
seed	optional; an object specifying if and how the random number generator should be initialized (‘seeded’). If seed is an integer, it is passed to <code>set.seed</code> prior to any simulation and is returned as the “seed” element of the return list. By default, the state of the random number generator is not changed and the value of <code>.Random.seed</code> on the call is stored in the “seed” element of the return list.
verbose	logical; if TRUE, progress information is reported as pfilter2 works.
...	By default, when <code>pfilter2 pfilter</code> is run on a <code>pfilterd2.pomp</code> object, the settings in the original call are re-used. This default behavior can be overridden by changing the settings (see Examples below).

Value

An object of class `pfilterd2.pomp`. This class inherits from class `pomp` and contains the following additional slots:

pred.mean, pred.var, filter.mean matrices of prediction means, variances, and filter means, respectively. In each of these, the rows correspond to states and parameters (if appropriate), in that order, the columns to successive observations in the time series contained in object.

eff.sample.size numeric vector containing the effective number of particles at each time point.

cond.loglik numeric vector containing the conditional log likelihoods at each time point.

saved.states If `pfilter2` was called with `save.states=TRUE`, this is the list of state-vectors at each time point, for each particle. It is a length-`ntimes` list of `nvars-by-Np` arrays. In particular, `saved.states[[t]][,i]` can be considered a sample from $f[X_t|y_{1:t}]$.

saved.params If `pfilter2` was called with `save.params=TRUE`, this is the list of parameter-vectors at each time point, for each particle. It is a length-`ntimes` list of `npars-by-Np` arrays. In particular, `saved.params[[t]][,i]` is the parameter portion of the i -th particle at time t .

seed the state of the random number generator at the time `pfilter2` was called. If the argument `seed` was specified, this is a copy; if not, this is the internal state of the random number generator at the time of call.

Np, tol, nfail the number of particles used, failure tolerance, and number of filtering failures, respectively.

loglik the estimated log-likelihood.

These can be accessed using the `$` operator as if the returned object were a list. In addition, `logLik` returns the log likelihood. Note that if the argument `params` is a named vector, then these parameters are included in the `params` slot of the returned `pfilterd2.pomp` object. That is `coef(pfilter2(obj,params=theta))==theta` if `theta` is a named vector of parameters.

Author(s)

Dao Nguyen <nguyenxd at umich dot edu>, Edward L. Ionides <ionides at umich dot edu>

References

M. S. Arulampalam, S. Maskell, N. Gordon, & T. Clapp. A Tutorial on Particle Filters for Online Nonlinear, Non-Gaussian Bayesian Tracking. IEEE Trans. Sig. Proc. 50:174–188, 2002.

See Also

[is2](#)

pfilter2-methods

Methods of the "pfilterd2.pomp" class

Description

Methods of the "pfilterd2.pomp" class.

Usage

```
## S4 method for signature pfilterd2.pomp
logLik(object, ...)
## S4 method for signature pfilterd2.pomp
pred.mean(object, pars, ...)
## S4 method for signature pfilterd2.pomp
pred.var(object, pars, ...)
## S4 method for signature pfilterd2.pomp
filter.mean(object, pars, ...)
## S4 method for signature pfilterd2.pomp
eff.sample.size(object, ...)
## S4 method for signature pfilterd2.pomp
cond.logLik(object, ...)
## S4 method for signature pfilterd2.pomp
as(object, class)
## S4 method for signature pfilterd2.pomp,data.frame
coerce(from, to = "data.frame", strict = TRUE)
## S3 method for class pfilterd2.pomp
as.data.frame(x, row.names, optional, ...)
```

Arguments

object, x	An object of class pfilterd2.pomp or inheriting class pfilterd2.pomp.
pars	Names of parameters.
class	character; name of the class to which object should be coerced.
from, to	the classes between which coercion should be performed.
strict	ignored.
row.names, optional	ignored.
...	Additional arguments unused at present.

Author(s)

Dao Nguyen <nguyenxd at umich dot edu>, Edward L. Ionides <ionides at umich dot edu>

See Also

[pfilter2](#)

Index

*Topic **models**
 is2-methods, 5
 pfilter2, 6
 pfilter2-methods, 8
*Topic **ts**
 is2, 2
 is2-methods, 5
 pfilter2, 6
 pfilter2-methods, 8
[, is2List-method (is2-methods), 5
[-is2List (is2-methods), 5
\$, pfilterd2.pomp-method
 (pfilter2-methods), 8
\$-pfilterd2.pomp (pfilter2-methods), 8

as, pfilterd2.pomp-method
 (pfilter2-methods), 8
as.data.frame.pfilterd2.pomp
 (pfilter2-methods), 8

c, is2-method (is2-methods), 5
c, is2List-method (is2-methods), 5
c-is2 (is2-methods), 5
c-is2List (is2-methods), 5
coef, 5
coerce, pfilterd2.pomp, data.frame-method
 (pfilter2-methods), 8
compare.is2 (is2-methods), 5
cond.logLik (pfilter2-methods), 8
cond.logLik, pfilterd2.pomp-method
 (pfilter2-methods), 8
cond.logLik-pfilterd2.pomp
 (pfilter2-methods), 8
continue (is2), 2
continue, is2-method (is2), 2
continue-is2 (is2), 2
conv.rec (is2-methods), 5
conv.rec, is2-method (is2-methods), 5
conv.rec, is2List-method (is2-methods), 5
conv.rec-is2 (is2-methods), 5

conv.rec-is2List (is2-methods), 5

eff.sample.size (pfilter2-methods), 8
eff.sample.size, pfilterd2.pomp-method
 (pfilter2-methods), 8
eff.sample.size-pfilterd2.pomp
 (pfilter2-methods), 8

filter.mean (pfilter2-methods), 8
filter.mean, pfilterd2.pomp-method
 (pfilter2-methods), 8
filter.mean-pfilterd2.pomp
 (pfilter2-methods), 8

is2, 2, 6, 8
is2, is2-method (is2), 2
is2, pfilterd2.pomp-method (is2), 2
is2, pomp-method (is2), 2
is2-class (is2), 2
is2-is2 (is2), 2
is2-methods, 5
is2-pfilterd2.pomp (is2), 2
is2-pomp (is2), 2
is2List-class (is2-methods), 5

logLik, is2-method (is2-methods), 5
logLik, pfilterd2.pomp-method
 (pfilter2-methods), 8
logLik-is2 (is2-methods), 5
logLik-pfilterd2.pomp
 (pfilter2-methods), 8

pfilter2, 3, 5, 6, 6, 9
pfilter2, pfilterd2.pomp-method
 (pfilter2), 6
pfilter2, pomp-method (pfilter2), 6
pfilter2-methods, 8
pfilter2-pfilterd2.pomp (pfilter2), 6
pfilter2-pomp (pfilter2), 6
pfilterd2.pomp, 7
pfilterd2.pomp-class (pfilter2), 6

`plot, is2-method` (`is2-methods`), 5
`plot, is2List-method` (`is2-methods`), 5
`plot-is2` (`is2-methods`), 5
`plot-is2List` (`is2-methods`), 5
`pomp`, 4, 7
`pred.mean` (`pfilter2-methods`), 8
`pred.mean, pfilterd2.pomp-method`
 (`pfilter2-methods`), 8
`pred.mean-pfilterd2.pomp`
 (`pfilter2-methods`), 8
`pred.var` (`pfilter2-methods`), 8
`pred.var, pfilterd2.pomp-method`
 (`pfilter2-methods`), 8
`pred.var-pfilterd2.pomp`
 (`pfilter2-methods`), 8